

# STAT 391 Homework 2

Chongyi Xu  
University of Washington  
STAT 391 Spring 2018  
chongyix@uw.edu

---

---

## 1. Problem 1 - Estimation of small probabilities

- a Estimate  $\hat{\theta}_a, \hat{\theta}_b, \dots, \hat{\theta}_z$  from the given text I basically use the same setup as I did for HW1 Question 4. The only difference is that I used  $\theta^M L = 0$  for those letters that never appear.

```
import math

# Helper Methods
def langReader(file):
    '''
    langReader is used to read in files and compute for the
    probability
    for each letter

    Args:
        file: The name of the file containing the letter and
        probability.

    Returns:
        A dictionary containing letters as keys and probability
        as values
    '''
    pLang = {}
    for line in open(file):
        el = line.split(' ')
        letter = el[1].lower()
        pLang[letter] = float(el[2])/1000
```

```

        return pLang

def LetterCounter(testString):
    '''
    Counts for the number of each letter in the test string
    (sufficient statistics)

    Args:
        testString: The string to count
    Returns:
        A dictionary containing letters as keys and count
        number as values
    '''
    counter = {}
    for char in testString:
        try:
            counter[char] = counter[char] + 1
        except KeyError:
            counter[char] = 1
    return counter

# Initialize the probability map
path =
    r'C:\Users\johnn\Documents\UW\SchoolWorks\2018Spring\STAT391\HW1'

english = langReader(path + r'\english.dat')

# Read in lincoln_text.txt
txt = ''
for line in open(path + r'\lincoln_text.txt'):
    txt = txt + line
txt = ''.join(e for e in txt if e.isalnum()).lower()
counter = LetterCounter(txt)
length = len(txt)

# Since there might be some missing letters
lincolnEng = {}
for letter in english:
    if letter not in counter:
        counter[letter] = 0

```

```
lincolnEng[letter] = counter[letter] / length
print(lincolnEng)
```

As the result, I got my  $\hat{\theta}$  to be

```
{'a': 0.07108239095315025, 'b': 0.015347334410339256,
'c': 0.020193861066235864, 'd': 0.029079159935379646,
'e': 0.11793214862681745, 'f': 0.02665589660743134,
'g': 0.021809369951534735, 'h': 0.05492730210016155,
'i': 0.07673667205169628, 'j': 0.0024232633279483036,
'k': 0.004038772213247173, 'l': 0.045234248788368334,
'm': 0.024232633279483037, 'n': 0.06704361873990307,
'o': 0.08239095315024232, 'p': 0.01615508885298869,
'q': 0.0008077544426494346, 'r': 0.08077544426494346,
's': 0.05977382875605816, 't': 0.0888529886914378,
'u': 0.03715670436187399, 'v': 0.01050080775444265,
'w': 0.018578352180936994, 'x': 0.0,
'y': 0.027463651050080775, 'z': 0.0008077544426494346}
```

- b Let  $S$  be the sample space  $\{a, b, c, \dots, z\}$ , with  $m = |S| = 26$ . Determine the sets  $S_0, S_1, \dots, S_n$ , where  $S_k = \{j \in S, n_j = k\}$ .

```
sample_count = LetterCounter(txt)
sk = {}
for letter in sample_count:
    if sample_count[letter] not in sk:
        sk[sample_count[letter]] = [letter]
    else:
        sk[sample_count[letter]].append(letter)
print(sk)
```

```
{23: ['w'], 68: ['h'], 88: ['a'], 110: ['t'], 25: ['c'],
102: ['o'], 83: ['n'], 74: ['s'], 95: ['i'], 46: ['u'],
146: ['e'], 19: ['b'], 56: ['l'], 100: ['r'], 5: ['k'],
33: ['f'], 34: ['y'], 36: ['d'], 20: ['p'], 27: ['g'],
30: ['m'], 13: ['v'], 1: ['z', 'q'], 3: ['j']}
```

- c Let  $r_k = |S_k|$  and  $r$  be the number of unique letters observed in the Lincoln-English corpus above. Verify that  $r = \sum_{k=1}^n r_k$ ,  $m = \sum k = 0^n r_k$ , and  $n = \sum_{k=1}^n k r_k$ .

```
r = sum(len(sk[k]) for k in sk)
print(r)

n = sum(k*len(sk[k]) for k in sk)
print(n)

print(len(txt))
```

And the result I got is

$$\begin{aligned}
 r &= |S_1| + |S_3| + \cdots + |S_{146}| \\
 &= 25 \\
 m &= |S_0| + r \\
 &= 26 \\
 n &= 1 * |S_1| + 3 * |S_3| + \cdots + 146 * |S_{146}| \\
 &= 1238 \\
 length &= 1238
 \end{aligned}$$

## 2. Problem 2 - Estimate letter probabilities from text

1. Get the sufficient statistics. Only print out  $n_{a:j}$

```
import math

# Helper Methods
def LetterCounter(testString):
    """
    Counts for the number of each letter in the test string
    (sufficient statistics)

    Args:
        testString: The string to count
    Returns:
```

```

        A dictionary containing letters as keys and count
        number as values
    '''
    counter = {}
    for char in testString:
        try:
            counter[char] = counter[char] + 1
        except KeyError:
            counter[char] = 1
    return counter

# Initialize the probability map
path =
    r'C:\Users\johnn\Documents\UW\SchoolWorks\2018Spring\STAT391\HW2'
english = langReader(path + r'\english.dat')

# Read in mlk-letter-estimation.txt
txt = ''
for line in open(path + r'\hw2-mlk-letter-estimation.txt'):
    txt = txt + line
txt = ''.join(e for e in txt if e.isalnum()).lower()
counter = LetterCounter(txt)
for letter in english: # Update missing letters
    if letter not in counter:
        counter[letter] = 0
length = len(txt)

for k in counter:
    if k <= 'j':
        print('\n'+k+'='+str(counter[k]))

```

And the output is

```

na=24
ne=44
nf=17
nh=19
ng=3
nd=10

```

```
ni=29
nc=18
nb=2
nj=0
```

What is the fingerprint  $r_k$  of this dataset?

```
sk = {}
for letter in counter:
    if counter[letter] not in sk:
        sk[counter[letter]] = [letter]
    else:
        sk[counter[letter]].append(letter)
rk = {}
for k in sk:
    rk[k] = len(sk[k])
    print('r'+str(k)+'='+str(rk[k]))
```

```
r39=1
r32=1
r16=1
r24=1
r4=1
r44=1
r11=1
r28=1
r17=1
r15=1
r19=1
r5=1
r3=2
r10=1
r29=1
r6=1
r18=1
r9=1
r2=2
r7=1
r0=4
```

2. Compute the ML estimates  $\theta_{A:Z}^{M,L}$

```
print('ML Estimation:')
ML = {}
for k in sk:
    letter = sk[k]
    for i in letter:
        ML[i] = k / length
        if letter.index(i) == 0:
            print('theta_'+i+'='+str(ML[i]))
        else:
            print('theta_'+i+'='+theta_+letter[0])
```

```
ML Estimation:
theta_t=0.11370262390670553
theta_o=0.09329446064139942
theta_s=0.04664723032069971
theta_a=0.06997084548104957
theta_v=0.011661807580174927
theta_e=0.1282798833819242
theta_m=0.03206997084548105
theta_n=0.08163265306122448
theta_f=0.04956268221574344
theta_r=0.043731778425655975
theta_h=0.05539358600583091
theta_p=0.014577259475218658
theta_g=0.008746355685131196
theta_w=theta_g
theta_d=0.029154518950437316
theta_i=0.08454810495626822
theta_y=0.01749271137026239
theta_c=0.052478134110787174
theta_u=0.026239067055393587
theta_b=0.0058309037900874635
theta_k=theta_b
theta_l=0.02040816326530612
theta_j=0.0
```

```

theta_q=theta_j
theta_x=theta_j
theta_z=theta_j

```

### 3. Compute Laplace $\theta_{A:Z}^{Lap}$ of the same probability

```

print('Laplace Estimation:')
lap = {}
m = 26
for k in sk:
    letter = sk[k]
    for i in letter:
        lap[i] = (k + 1) / (length + m)
        if letter.index(i) == 0:
            print('theta_'+i+'='+str(lap[i]))
        else:
            print('theta_'+i+'='+theta_+letter[0])

```

```

Laplace Estimation:
theta_t=0.10840108401084012
theta_o=0.08943089430894309
theta_s=0.04607046070460705
theta_a=0.06775067750677506
theta_v=0.013550135501355014
theta_e=0.12195121951219512
theta_m=0.032520325203252036
theta_n=0.07859078590785908
theta_f=0.04878048780487805
theta_r=0.04336043360433604
theta_h=0.05420054200542006
theta_p=0.016260162601626018
theta_g=0.01084010840108401
theta_w=theta_g
theta_d=0.02981029810298103
theta_i=0.08130081300813008
theta_y=0.018970189701897018
theta_c=0.051490514905149054
theta_u=0.02710027100271003

```



```

theta_b=0.008130081300813009
theta_k=theta_b
theta_l=0.02168021680216802
theta_j=0.0027100271002710027
theta_q=theta_j
theta_x=theta_j
theta_z=theta_j

```

4. Compute the Witten-Bell estimates  $\theta_{A:Z}^{W,B}$

```

r = sum(len(sk[k]) for k in sk if k!=0)
print('Witten-Bell Estimation:')
wb = {}
for k in sk:
    letter = sk[k]
    for i in letter:
        if k != 0:
            wb[i] = k / (length + r)
        else:
            wb[i] = 1 / (m - r) * r / (length + r)
        if letter.index(i) == 0:
            print('theta_'+i+'='+str(wb[i]))
        else:
            print('theta_'+i+'='+theta_+letter[0])

```

```

Witten-Bell Estimation:
theta_t=0.10684931506849316
theta_o=0.08767123287671233
theta_s=0.043835616438356165
theta_a=0.06575342465753424
theta_v=0.010958904109589041
theta_e=0.12054794520547946
theta_m=0.030136986301369864
theta_n=0.07671232876712329
theta_f=0.04657534246575343
theta_r=0.0410958904109589
theta_h=0.052054794520547946
theta_p=0.0136986301369863

```

```

theta_g=0.00821917808219178
theta_w=theta_g
theta_d=0.0273972602739726
theta_i=0.07945205479452055
theta_y=0.01643835616438356
theta_c=0.049315068493150684
theta_u=0.024657534246575342
theta_b=0.005479452054794521
theta_k=theta_b
theta_l=0.019178082191780823
theta_j=0.015068493150684932
theta_q=theta_j
theta_x=theta_j
theta_z=theta_j

```

5. Compute the smoothed Good-Turing estimates  $\theta_{A:Z}^{GT}$

```

print('Good-Turing Estimation:')
gt = {}
for k in sk:
    letter = sk[k]
    for i in letter:
        try:
            gt[i] = ((k + 1) * rk[k+1] / rk[k]) / length
        except KeyError:
            gt[i] = 0
    if letter.index(i) == 0:
        print('theta_'+i+'='+str(gt[i]))
    else:
        print('theta_'+i+'='+theta_+letter[0])

```

```

Good-Turing Estimation:
theta_t=0
theta_o=0
theta_s=0.04956268221574344
theta_a=0
theta_v=0.014577259475218658
theta_e=0

```

```

theta_m=0
theta_n=0.08454810495626822
theta_f=0.052478134110787174
theta_r=0.04664723032069971
theta_h=0
theta_p=0.01749271137026239
theta_g=0.0058309037900874635
theta_w=theta_g
theta_d=0.03206997084548105
theta_i=0
theta_y=0.02040816326530612
theta_c=0.05539358600583091
theta_u=0.029154518950437316
theta_b=0.008746355685131196
theta_k=theta_b
theta_l=0
theta_j=0
theta_q=theta_j
theta_x=theta_j
theta_z=theta_j

```

6. Compute the Ney-Essen estimates  $\theta_{A:Z}^{N,E}$ , taking  $\delta = 1$

```

print('\nNey-Essen Estimation:')
D = 0
delta = 1
for letter in counter:
    D = D + min(counter[letter], delta)
ne = {}
for k in sk:
    letter = sk[k]
    for i in letter:
        ne[i] = (k - min(k, delta) + D / m) / length
        if letter.index(i) == 0:
            print('theta_'+i+'='+str(ne[i]))
        else:
            print('theta_'+i+'='+theta_+letter[0])

```

```

Ney-Essen Estimation:
theta_t=0.11325409284592958
theta_o=0.09284592958062346
theta_s=0.04619869925992375
theta_a=0.06952231442027361
theta_v=0.01121327651939897
theta_e=0.12783135232114823
theta_m=0.031621439784705094
theta_n=0.08118412200044853
theta_f=0.049114151154967485
theta_r=0.04328324736488002
theta_h=0.054945054945054944
theta_p=0.014128728414442699
theta_g=0.008297824624355236
theta_w=theta_g
theta_d=0.02870598788966136
theta_i=0.08409957389549226
theta_y=0.017044180309486432
theta_c=0.05202960305001122
theta_u=0.025790535994617628
theta_b=0.005382372729311505
theta_k=theta_b
theta_l=0.01995963220453016
theta_j=0.002466920834267773
theta_q=theta_j
theta_x=theta_j
theta_z=theta_j

```

7. Now use the estimates you obtained to compute the log-probability of the text in hw2-test-letter-estimation-large.txt. Also compute the log-probability of the training data hw2-mlk-letter-estimation.txt

```

from decimal import Decimal
# Helper Methods
def computeP(counter, probMap):
    '''
    Compute for the probability of with given letter counter
    and language probability map

```

```

    Args:
        counter: The letter counter of the test string.
        probMap: The probability map of the test language.
    Returns:
        A integer telling P(sentence)
    """
    p = 1
    for letter in counter:
        try:
            p = p * Decimal(probMap[letter]**counter[letter])
        except KeyError:
            print("The letter ", letter, " is not in this
                  language")
    return p

def MaxLogLikelihood(fileName, langDict):
    """
    Find the maximum log-likelihood according to the fileName
    and output the guess.

    Args:
        fileName: The file name with txt to test
        langDict: The dictionary for all languages.
    Output:
        The log-likelihood for each language and the best guess.
    """
    print("\nConsidering the file: ", fileName)
    # Remove spaces and punctuation
    txt = ''
    for line in open(path + fileName):
        txt = txt + line
    testStr = ''.join(e for e in txt if e.isalnum()).lower()
    counter = LetterCounter(testStr)
    best = [-float('inf'), '']
    for lang in langDict:
        p = computeP(counter, langDict[lang])
        ll = p.ln() / Decimal(math.log(2))
        print("The log-likelihood for ", lang, " is ", ll)
        if best[0] < ll:

```

```

        best = [ll, lang]
        print("And as the result, the best guess is ", best[1], "
              with likelihood ", best[0], "\n")

langDict = {'ML Estimation':ML,\
            'Laplace Estimation':lap,\
            'Witten-Bell Estimation':wb,\
            'Smoothed Good-Turning Estimation':gt,\
            'Ney-Essen Estimation':ne}

MaxLogLikelihood(r'\hw2-mlk-letter-estimation.txt', langDict)
MaxLogLikelihood(r'\hw2-test-letter-estimation-large.txt',
                 langDict)

```

Considering the file: \hw2-mlk-letter-estimation.txt  
The log-likelihood for ML Estimation is  
-1380.953521626853747487385781  
The log-likelihood for Laplace Estimation is  
-1387.341205133273515356164497  
The log-likelihood for Witten-Bell Estimation is  
-1411.716467071790303072722354  
The log-likelihood for Smoothed Good-Turning Estimation is  
-Infinity  
The log-likelihood for Ney-Essen Estimation is  
-1385.893096046358672696316814  
And as the result, the best guess is ML Estimation with likelihood  
-1380.953521626853747487385781

Considering the file: \hw2-test-letter-estimation-large.txt  
The log-likelihood for ML Estimation is  
-Infinity  
The log-likelihood for Laplace Estimation is  
-8814.732094912244450343732977  
The log-likelihood for Witten-Bell Estimation is  
-8969.144788769402325042405529  
The log-likelihood for Smoothed Good-Turning Estimation is  
-Infinity

The log-likelihood for Ney-Essen Estimation is  
-8835.935483824767400464608135  
And as the result, the best guess is Laplace Estimation with  
likelihood  
-8814.732094912244450343732977

From the result above, we can see that Laplace Estimation gives the highest likelihood for the new text data. And Max-Likelihood Method gives the highest likelihood for the training data.

### 3. Problem 3 - ML estimation

1. Write the expression of the probability  $P(3, 2, 1, 1, 6)$

$$P(3, 2, 1, 1, 6) = \theta_o^3 \theta_e$$

2. Write the expression of  $l(\theta_o, \theta_e)$  the log-likelihood of data set  $D$  as a function of  $\theta_o, \theta_e$  and the counts  $n_{1:6}$

$$l(\theta_o, \theta_e) = n_1 \log \theta_o + n_2 \log \theta_e + n_3 \log \theta_o + n_4 \log \theta_e + n_5 \log \theta_o + n_6 \log \theta_e$$

3. Transform  $l(\theta_o, \theta_e)$  into a function of one variable  $l(\theta_e)$

$$l(\theta_o, \theta_e) = n_1 \log\left(\frac{1}{3} - \theta_e\right) + n_2 \log \theta_e + n_3 \log\left(\frac{1}{3} - \theta_e\right) + n_4 \log \theta_e + n_5 \log\left(\frac{1}{3} - \theta_e\right) + n_6 \log \theta_e$$

4. Now find the ML estimate of  $\theta_e$  by equating the derivative of  $l(\theta_e)$  with 0.

$$\begin{aligned} l'(\theta_e) = 0 &= -\frac{n_1}{\frac{1}{3} - \theta_e} + \frac{n_2}{\theta_e} - \frac{n_3}{\frac{1}{3} - \theta_e} + \frac{n_4}{\theta_e} - \frac{n_5}{\frac{1}{3} - \theta_e} + \frac{n_6}{\theta_e} \\ 0 &= \frac{-3(n_1 + n_3 + n_5)}{1 - 3\theta_e} + \frac{n_2 + n_4 + n_6}{\theta_e} \\ 0 &= -3(n_1 + n_3 + n_5)\theta_e + (n_2 + n_4 + n_6)(1 - 3\theta_e) \\ n_2 + n_4 + n_6 &= 3\theta_e(n_1 + \dots + n_6) \\ \theta_e &= \frac{n_2 + n_4 + n_6}{3 \sum_{i=1}^6 n_i} \end{aligned}$$

5. Explain why this result is natural

This result is reasonable since the  $\theta_e$  is found by the number of evens divided by 3, which is exactly the result I got from the computation above. Therefore, the ML I found is intuitive.

#### 4. Problem4 - The ML estimate as a random variable

1. What is the set of possible values  $S_{\theta_1}$  for  $\theta_1^{ML}$ ? Does the true  $\theta_1$  belong to  $S_{\theta_1}$ ?

The set of possible values  $S_{\theta_1}$  is

$$S_{\theta_1} = \left\{ \frac{100}{i} \text{ for } i \in \mathbb{N}, 0 \leq i \leq 100 \right\}$$

The true  $\theta_1$  does not belong to this set since  $i = 100 * 0.3141 = 31.41 \notin \mathbb{N}$

2. Write the expression of the probability of each outcome in  $S_{\theta_1}$

$$P(\theta_1^{ML} = \frac{i}{100}) = \binom{i}{100} \theta_1^i (1 - \theta_1)^{100-i}$$

3. Make a plot of the probability distribution of  $\theta_1^{ML}$

```
import numpy as np
import math
import matplotlib.pyplot as plt

n = 100
p = 0.3141

log_P = [0]*101
for i in range(0, 101): # using ln-gamma to avoid overflow
    log_P[i] = math.log(math.gamma(n+1)) -
               math.log(math.gamma(i+1)) - \
               math.log(math.gamma(n-i+1)) + i*math.log(p) +
               (n-i)*math.log(1-p)
theta = [math.exp(log_P[i]) for i in range(len(log_P))]
```



```

x = np.arange(0, 1.01, step=0.01)
plt.stem(x, theta)
plt.xlabel('theta_1')
plt.title('Probability of theta_1')
plt.show()

plt.stem(x, theta)
plt.xlabel('theta_1')
plt.title('Probability of theta_1')
plt.xlim([0.195, 0.405])
plt.show()

```

Figure 1 is the probability distribution of  $\theta_1$ .

4. Let  $\epsilon = 0.02$ . Answer using the probability distribution computed previously

```

e = 0.02
print('P{Absolute Error > 0.02} =', \
      sum(theta[i] for i in np.where(abs(x - p) > e)[0]))
print('P{Related Error > 0.02} =', \
      sum(theta[i] for i in np.where(abs(((x - p) / p)) > e)[0]))

```

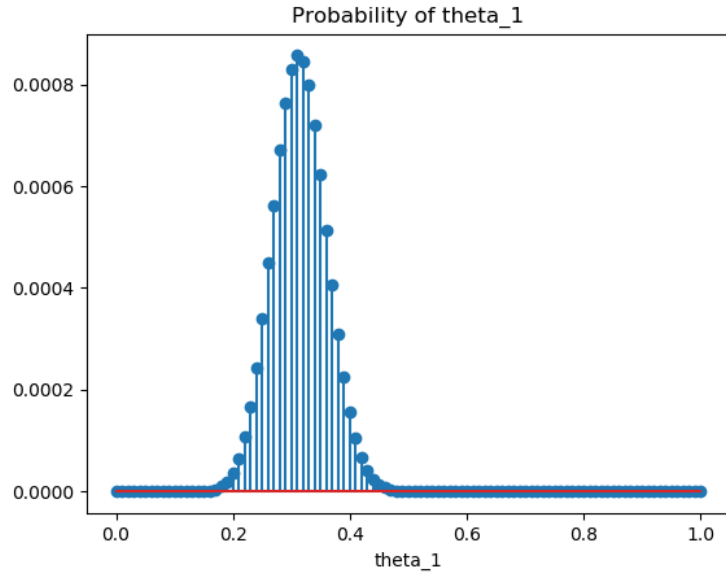
$$\begin{aligned}\delta_{abs} &= P[|\theta_1^{ML} - \theta_1| > \epsilon] &&= 0.6671038038886307 \\ \delta_{rel} &= P\left[\left|\frac{\theta_1^{ML} - \theta_1}{\theta_1}\right| > \epsilon\right] &&= 0.829757447824191\end{aligned}$$

5. For  $\epsilon = 0, 0.005, 0.001, \dots, 1$  plot the graph of  $\delta_{abs} = P[|\theta_1^{ML} - \theta_1| > \epsilon]$  vs  $\epsilon$ . Is the function  $\delta(\epsilon)$  monotonically increasing, decreasing or neither?

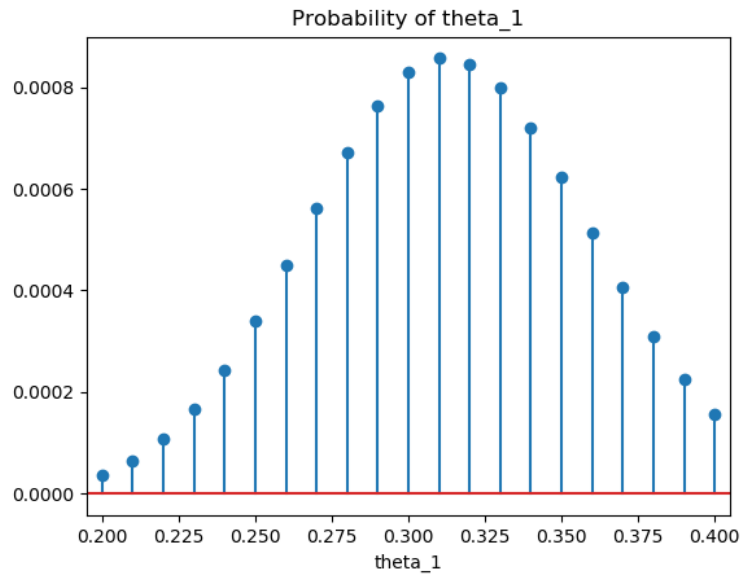
```

i = 0
epsilons = np.arange(0, 1, step=0.005)
delta = [0]*len(epsilons)
for e in epsilons:
    delta[i] = sum(theta[i] for i in np.where(abs(x - p) > e)[0])
    i = i + 1

```



(a) Plot of probability distribution of  $\theta_1$



(b) Plot of probability distribution of  $\theta_1$ , enlarged from  $\theta_1 = 20$  to  $\theta_1 = 40$

Figure 1: Probability distribution of  $\theta_1$

```
plt.plot(epsilons, delta)
plt.xlabel('epsilon')
plt.title('Delta(epsilon) vs. epsilon')
plt.show()
```

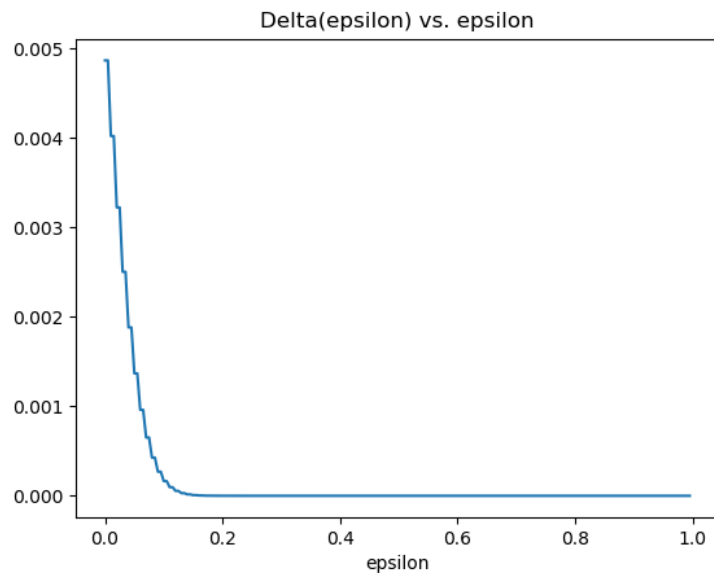


Figure 2:  $\delta(\epsilon)$  vs.  $\epsilon$

From the figure 2, we can clearly see that the function  $\delta(\epsilon)$  is monotonically decreasing(non-increasing).

6. Simulate tossing the coin with  $\theta_1 = 0.3141$ ,  $n = 100$  times and compute  $\theta_1^{ML}$ . What is the value you  $\theta_1^{ML}$  have obtained, and what are the absolute and relative error?

```
# Simulation
np.random.seed(999)
y = sum(np.random.binomial(1, p, n))
theta = y / n
print(theta - p)
print((theta - p) / p)
```

Absolute Error of Simulation = 0.085900000000000003

Relative Error of Simulation = 0.27347978350843694

Further, I have made 1000 observations to see the behavior of simulated  $\theta_1$

```
observations = 1000

yy = np.repeat(0, n+1)
for k in range(observations):
    yi = sum(np.random.binomial(1, p, n))
    yy[yi] = yy[yi] + 1

theta = [yy[i] / (n*observations) for i in range(0, n+1)]
x = np.arange(0, 1.01, step=0.01)
plt.stem(x, theta)
plt.xlabel('Simulated theta_1')
plt.title('Probability distribution of simulated theta_1')
plt.show()

e = 0.02
print('P{Absolute Error > 0.02}=', \
      sum(theta[i] for i in np.where(abs(x - p) > e)[0]))
print('P{Related Error > 0.02}=', \
      sum(theta[i] for i in np.where(abs((x - p) / p) >
      e)[0]))
```

P{Absolute Error > 0.02} = 0.67300000000000002

P{Related Error > 0.02} = 0.81200000000000003

7. Let  $\theta'_1$  have the value  $\theta_1^{ML}$  of the previous question. Repeat Question 3-6 using "the guess  $\theta'_1$  instead of "the truth"  $\theta$ .

```
#=====#
# Use the guess theta
p = theta_1

log_P = [0]*101
for i in range(0, 101): # using ln-gamma to avoid overflow
```

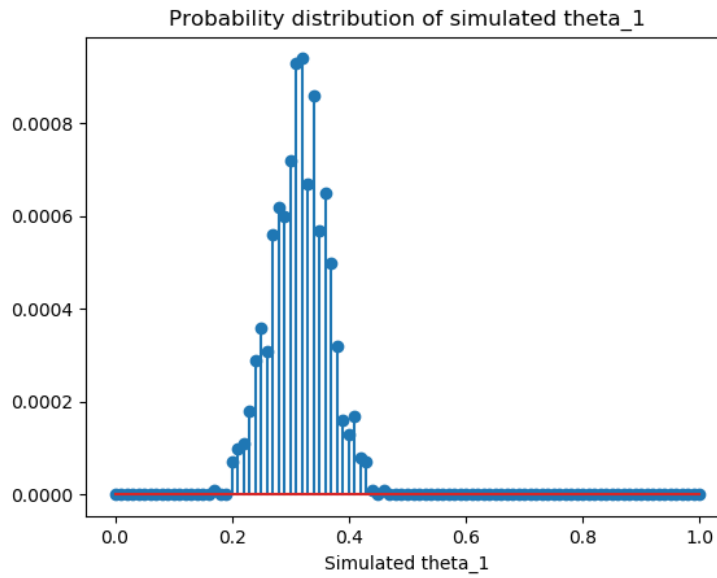


Figure 3: Probability distribution of simulated  $\theta_1$

```

log_P[i] = math.log(math.gamma(n+1)) -
    math.log(math.gamma(i+1)) -\
    math.log(math.gamma(n-i+1)) + i*math.log(p) +
    (n-i)*math.log(1-p)
y = [math.exp(log_P[i]) for i in range(len(log_P))]
plt.stem(y)
plt.xlabel('k')
plt.title('Probability of observing outcome 1 with k times')
plt.show()

theta = [y[i] / n for i in range(0, n+1)]
x = np.arange(0, 1.01, step=0.01)
plt.stem(x, theta)
plt.xlabel('theta_1')
plt.title('Probability of theta_1')
plt.show()

e = 0.02
print('P{Absolute Error > 0.02}=',\
    sum(theta[i] for i in np.where(abs(x - p) > e)[0]))

```

```

print('P{Related Error > 0.02}= ', \
      sum(theta[i] for i in np.where(abs((x - p) / p) >
                                      e)[0]))

i = 0
epsilons = np.arange(0, 1, step=0.005)
delta = [0]*len(epsilons)
for e in epsilons:
    delta[i] = sum(theta[i] for i in np.where(abs(x - p) >
                                                e)[0])
    i = i + 1
plt.plot(epsilons, delta)
plt.xlabel('epsilon')
plt.title('Delta(epsilon) vs. epsilon')
plt.show()

# Simulation
np.random.seed(999)
y = sum(np.random.binomial(1, p, n))
theta = y / n
print('Absolute Error of Simulation = ', abs(theta - p))
print('Relative Error of Simulation = ', abs((theta - p) / p))

```

From the figure 6, we can see that the function is still monotonically decreasing(non-increasing). And for the error part,

```

P{Absolute Error > 0.02}= 0.00685447787126666
P{Related Error > 0.02}= 0.009187808550038788
Absolute Error of Simulation = 0.06999999999999995
Relative Error of Simulation = 0.17499999999999988

```

## 5. Problem 5- Rare outcomes and data set size

1. What is the probability that the outcome sequence contains no 1's?

$$P(n_1 = 0) = \theta_1^0 * (1 - \theta_1)^n$$

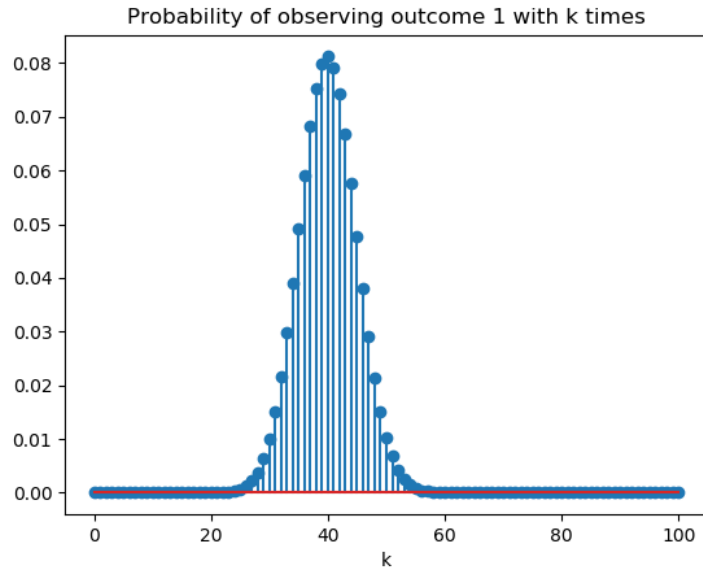


Figure 4: Probability distribution of outcome 1 using guess  $\theta$

2. What is the probability that the outcome sequence contains a single 1?

$$P(n_1 = 1) = \theta_1^1 * (1 - \theta_1)^{n-1}$$

3. Solve  $p_0 = p_1$

$$\begin{aligned} p_0 &\approx p_1 \\ P(n_1 = 0) &= P(n_1 = 0) + \epsilon \\ \theta_1^0 * (1 - \theta_1)^n &= \theta_1^1 * (1 - \theta_1)^{n-1} + \epsilon \\ (1 - \theta_1)^n &= \theta_1^1 * (1 - \theta_1)^{n-1} + \epsilon \end{aligned}$$

Say, the tolerance to be  $\epsilon = 10^{-10}$

4. Compute the above n for  $\theta_1 = 10^{-3}, 10^{-4}, 10^{-5}$ .

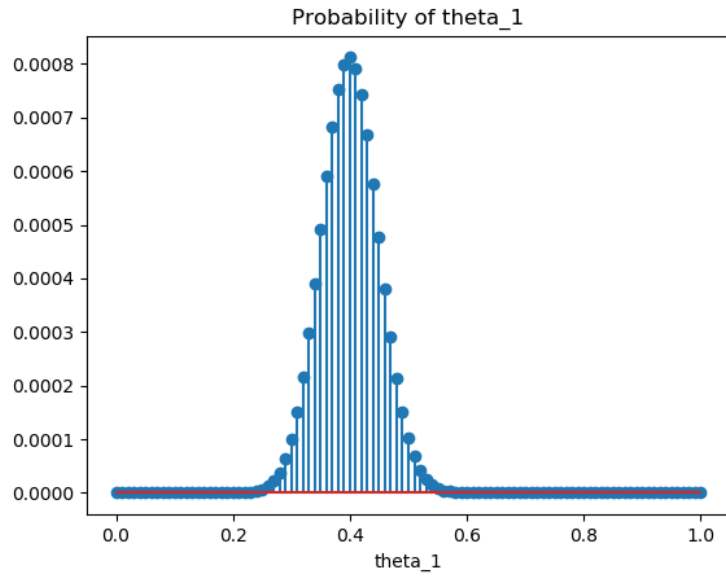


Figure 5: Probability distribution of  $\theta_1$  using guess  $\theta$

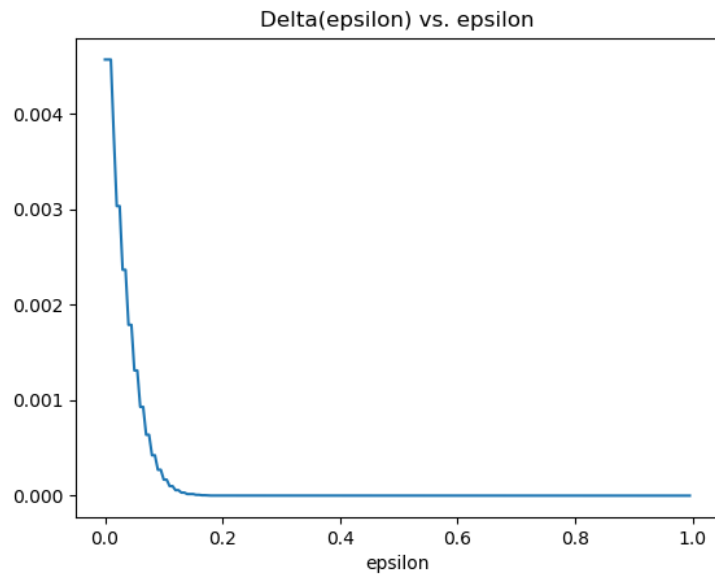


Figure 6:  $\delta(\epsilon)$  vs.  $\epsilon$  using guess  $\theta$